

Principles of Database Systems



Intermediate SQL (2)



主讲教师:

薛昕惟


办公室: 信息楼320B

xuexinwei@dlut.edu.cn

Constituent Parts of SQL (SQL组成部分)



- The SQL language has several parts:
 - Data-definition language (DDL)
 - Data-manipulation language (DML)
 - **Integrity** (完整性) (**included in DDL**)
 - **View definition** (视图定义) (**included in DDL**)
 - **Transaction control** (事务控制)
 - Authorization (授权)
 - Embedded SQL and dynamic SQL (嵌入式SQL及动态SQL)



Views

Views



- In some cases, it is not desirable for all users to see the entire logical model.
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

select ID, name, dept_name
from instructor

- **Any disadvantages?**



Views



- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “**virtual relation**” is called a **view**. (不是概念模型的一部分，对用户可见的“虚关系”)

Views



- SQL provides the **create view** command to create a view, which is stored in the database **in the long run**
 - **create view** v **as** <query expression>
 - where
 - <query expression> is any legal expression
 - the **view** name is represented by v
- View definition is not the same as creating a new relation by evaluating the query expression. (创建视图与创建关系不同)
 - Rather, a view definition causes the **saving of an expression**; the expression is substituted into queries using the view. (存储的是表达式)

Try...

- A view of instructors without their salary


```
create view faculty as  
select ID, name, dept_name  
from instructor
```



Try...

- Find all instructors in the Biology department

```
select name  
from faculty  
where dept_name = 'Biology'
```



Try...



- Create a view of department salary totals

```
create view  
departments_total_salary(dept_name, total_salary)  
as  
    select dept_name, sum (salary)  
    from instructor  
    group by dept_name;
```

Views Defined Using Other Views



- Explain the following views:
- **create view** `physics_fall_2009` **as**
 select course.course_id, sec_id, building, room_number
 from course, section
 where course.course_id = section.course_id
 and course.dept_name = 'Physics'
 and section.semester = 'Fall'
 and section.year = '2009';
- **create view** `physics_fall_2009_watson` **as**
 select course_id, room_number
 from `physics_fall_2009`
 where building= 'Watson';

Views Defined Using Other Views



- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as  
(select course_id, room_number  
from (select course.course_id, building,  
room_number  
      from course, section  
      where course.course_id = section.course_id  
           and course.dept_name = 'Physics'  
           and section.semester = 'Fall'  
           and section.year = '2009')  
where building= 'Watson');
```

Materialized Views



- **Materializing a view**(物化视图): **create a physical table** containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
 - Need to **maintain** the view(维护视图), by updating the view whenever the underlying relations are updated.

Drop View



- The **Drop View** command deletes the definition the view from the **data dictionary**.

drop view view_name;

Other views depending on this dropped view should be deleted explicitly.

Update of a View



- Add a new tuple to faculty view which we defined earlier
insert into faculty **values** ('30765', 'Green', 'Music');
- Two reasonable approaches:
 - Reject the insertion
 - Insert the tuple
('30765', 'Green', 'Music', null)
into the instructor relation
(必须转化为对实际关系的修改)

Some Updates cannot be Translated Uniquely




- **create view** instructor_info **as**
 select ID, name, building
 from instructor, department
 where instructor.dept_name=
department.dept_name;
- **insert into** instructor_info **values**
('69987', 'White', 'Taylor');
 - which department, if multiple departments in Taylor?
 - what if no department is in Taylor?

Some Updates cannot be Translated Uniquely



- Most SQL implementations allow updates only on simple views
 - The **from** clause has only **one** database relation.
 - The **select** clause contains only attribute names of the relation, and does **not have any expressions, aggregates, or distinct** specification.
 - Any attribute not listed in the **select** clause can be set to null.
 - The query does **not have** a **group** by or **having** clause.

And Some Not at All



- **create view** history_instructors **as**
 select *
 from instructor
 where dept_name= 'History';
- What happens if we insert
 ('25566', 'Brown', 'Biology', 100000) into
 history_instructors?
- **with check option:** if a tuple inserted into the
 view does not satisfy the view's **where clause**
 condition, the insertion is rejected by the database
 system



Transactions

Transactions



- A transaction is a sequence of queries and update statements on DB, executed as a single, and are started implicitly and terminated by one of commit work(提交)or rollback/abort work
 - **Commit work** : makes the updates performed by the transaction become permanent in the database.
 - **Rollback work**: undoes all the updates performed by the SQL statements in the transaction.

Transactions



- Unit of work
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Transactions begin implicitly
 - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
 - In SQL:1999, can use: **begin atomic end**
 - Not supported on most databases



Integrity Constraints

Integrity Constraints




- Integrity constraints guard **against accidental damage** to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - An instructor name cannot be *null*.
 - No two instructors can have the same instructor ID.
 - The budget of a department must be greater than \$0.00.

Integrity Constraints on a Single Relation



- primary key
- not null
- Unique
- foreign keys
- check (P), where P is a predicate

Not Null



- Declare name and budget to be **not null**
 - name **varchar(20) not null**
 - budget **numeric(12,2) not null**

Unique Constraints



- **unique** (A_1, A_2, \dots, A_m)
 - The unique specification states that the attributes A_1, A_2, \dots, A_m form a **candidate key**.
 - Candidate keys are **permitted to be null** (in contrast to primary keys).

The check clause



- The **check** clause is applied to relation declaration
 - check (P), where P is a predicate which must be satisfied by every tuple in the relation.
- Example: ensure that the budget of a department must be greater than \$0.00
 - **create table** department
(dept name **varchar (20)**,
building **varchar (15)**,
budget **numeric (12,2)**,
primary key (dept name)
check(budget>0));

Try...



- Ensures that semester is one of fall, winter, spring or summer:

```
create table section (  
  course_id varchar (8),  
  sec_id varchar (8),  
  semester varchar (6),  
  year numeric (4,0),  
  building varchar (15),  
  room_number varchar (7),  
  time slot id varchar (4),  
  primary key (course_id, sec_id, semester, year),  
  check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```

Referential Integrity



- Ensures that a value that appears in one relation for a given set of attributes **also appears** for a certain set of attributes in another relation.
 - Example: If “Biology” is a department name appearing in one of the tuples in the *course* relation, then there exists a tuple in the *department* relation for “Biology”. -- **foreign key**


```
create table course (  
    course_id char(5) primary key,  
    title varchar(20),  
    dept_name varchar(20) references department  
)
```

Cascading Actions in Referential Integrity



- When the DB is modified by **Insert**, **Delete**, and **Update**, the tests must be made in order to preserve the referential integrity constraint.
- **create table** course (
 ...
 dept_name **varchar**(20),
 foreign key (dept_name) **references** department
 on delete cascade
 on update cascade,
 ...
)
- alternative actions to **cascade**: **set null**, **set default**

Integrity Constraint Violation During Transactions



- E.g. **create table** person (
ID **char**(10) **primary key** ,
name **char**(40),
spouse **char**(10),
foreign key spouse **references** person)
- How to insert a tuple without causing constraint violation ?
 - set spouse to null initially, update after inserting all persons (not possible if spouse attribute declared to be **not null**)
 - OR **defer**(延迟) constraint checking

Defer Constraint Checking



- The SQL standard allows a clause **initially deferred** to be added to a constraint specification.
- For constraints declared as **deferrable**, executing a statement **set constraints** constraint-list **deferred** as part of a transaction causes the checking of the specified constraints to be deferred to the end of that transaction.

Complex Check Conditions and Assertions



- **check** (time_slot_id
 in (**select** time_slot_id **from** time_slot))
- Unfortunately: subquery in check clause not supported by pretty much any database.

Assertions



- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
 - e.g. domain constraints, referential-integrity constraint
- An assertion in SQL takes the form
create assertion <assertion-name> **check**
<predicate>

Assertions




- E.g. The value of the attribute tot_cred for each student must equal the sum of credits of courses that the student has completed successfully.

```
create assertion credits_earned_constraint check  
  (not exists (select ID from student  
    where tot_cred < > (  
      select sum(credits)  
      from takes join course  
      on takes.course_id= course.course_id  
      where student.ID=takes.ID and grade is not  
        null and grade < > 'F')
```

Assertions



- When an assertion is made, the DBMS tests it for validity. Any modification to DB is allowed only if it does not cause that assertion to be violated.
 - This testing may introduce a significant amount of overhead, hence **assertions should be used with great care.**
- Not supported by every DBMS.

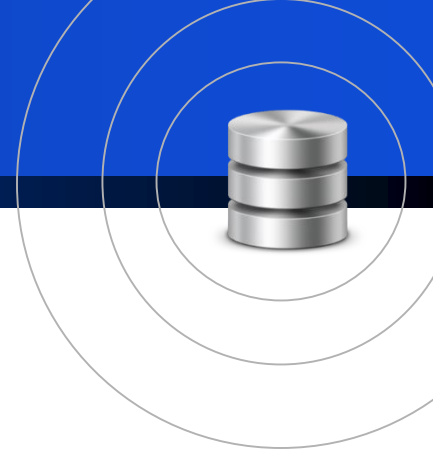


Review

Review



- Join Expressions
 - left outer join, right outer join, full outer join
 - inner join = join
 - Join types and join conditions
- Views
 - Create view
 - Use views in SQL queries
 - Update view: with check option
- Transactions
 - Atomic
 - Commit work, Rollback work



- Integrity Constraints
 - **Not null**
 - **unique** (A_1, A_2, \dots, A_m), candidate key, null
 - **Check**(P)
 - Referential Integrity, foreign key, **on delete/update cascade**, on delete/update set null, on delete/update set default
 - defer constraint checking
 - **create assertion** <assertion-name> **check** <predicate>